

cats_and_dogs

March 4, 2023

```
[1]: import tensorflow as tf
      %matplotlib inline
      import matplotlib.pyplot as plt
      plt.style.use('dark_background')
      import numpy as np
      import tensorflow_datasets as tfds
      import tensorflow_addons as tfa
      import os
      import random
```

```
[2]: (train_ds, validation_ds, test_ds), metadata = tfds.load('cats_vs_dogs',
                                                             split=['train[:80%]',
                                                             ↪ 'train[80%:90%]', 'train[90%:]'],
                                                             with_info=True,
                                                             as_supervised=True)
```

Metal device set to: Apple M1 Pro

systemMemory: 16.00 GB

maxCacheSize: 5.33 GB

2023-03-04 15:26:57.371182: I

tensorflow/core/common_runtime/pluggable_device/pluggable_device_factory.cc:305]
Could not identify NUMA node of platform GPU ID 0, defaulting to 0. Your kernel
may not have been built with NUMA support.

2023-03-04 15:26:57.371320: I

tensorflow/core/common_runtime/pluggable_device/pluggable_device_factory.cc:271]
Created TensorFlow device (/job:localhost/replica:0/task:0/device:GPU:0 with 0
MB memory) -> physical PluggableDevice (device: 0, name: METAL, pci bus id:
<undefined>)


```
[3]: dataset = train_ds.map(
      lambda image, label: (tf.image.convert_image_dtype(image, tf.float64),
      ↪ label)
      ).take(5)
```

```
[4]: f, axarr = plt.subplots(1,5)

i = 0
for image, label in dataset:
    print('Image shape: ', np.shape(image))
    tf.print('Label: ', label)
    axarr[i].imshow(image)
    axarr[i].axis('off')
    i += 1
```

```
Image shape: (262, 350, 3)
Label: 1
Image shape: (409, 336, 3)
Label: 1
Image shape: (493, 500, 3)
Label: 1
Image shape: (375, 500, 3)
Label: 0
Image shape: (240, 320, 3)
Label: 1
```

```
2023-03-04 15:27:03.271220: W
tensorflow/core/platform/profile_utils/cpu_utils.cc:128] Failed to get CPU
frequency: 0 Hz
```



```
[5]: rnd = np.random.RandomState(seed=1234)
# we need to convert images to same shape
# we need to convert images to same shape
IMG_HEIGHT = 224
IMG_WIDTH = 224

train = train_ds.map(
    lambda image, label: (tf.image.convert_image_dtype(image, tf.float64),
        ↪label)
).map( # resize for input
    lambda image, label: (tf.image.resize(image, [IMG_HEIGHT, IMG_WIDTH]),
        ↪label)
).map( # random horizontal flip
```

```

        lambda image, label: (tf.image.random_flip_left_right(image, seed=rnd.
↪seed(1234))), label)
    ).map( # random translation
        lambda image, label: (tfa.image.translate(image, tf.random.
↪normal(shape=[2], seed=1234))), label)
    ).take(5)

```

```

[6]: f, axarr = plt.subplots(1,5)

i = 0
for image, label in train:
    print('Image shape: ', np.shape(image))
    tf.print('Label:',label)
    axarr[i].imshow(image)
    axarr[i].axis('off')
    i += 1

```

```

Image shape: (224, 224, 3)
Label: 1
Image shape: (224, 224, 3)
Label: 1
Image shape: (224, 224, 3)
Label: 1
Image shape: (224, 224, 3)
Label: 0
Image shape: (224, 224, 3)
Label: 1

```



```

[75]: rnd = np.random.RandomState(seed=None)
num_threads = 4

# putting all the data processing pipeline on the CPU
# to make sure that the GPU is only used for training the deep neural network ↵
↪model
with tf.device('/cpu:0'):
    train = train_ds.map(
        lambda image, label: (tf.image.convert_image_dtype(image, tf.float64), ↵
↪label)

```

```

        ).map( # resize for input
            lambda image, label: (tf.image.resize(image, [IMG_HEIGHT, IMG_WIDTH]),
        ↪label)
        ).map( # random horizontal flip
            lambda image, label: (tf.image.random_flip_left_right(image, seed=rnd.
        ↪seed(1234)), label), num_parallel_calls=num_threads
        ).map( # random translation
            lambda image, label: (tfa.image.translate(image, tf.random.
        ↪normal(shape=[2], seed=rnd.seed(1234))), label),
        ↪num_parallel_calls=num_threads
        ).shuffle(1000).batch(32).take(1).prefetch(1)

with tf.device('/cpu:0'):
    valid = validation_ds.map(
        lambda image, label: (tf.image.convert_image_dtype(image, tf.
        ↪float64), label)
        ).map( # resize for input
            lambda image, label: (tf.image.resize(image, [IMG_HEIGHT,
        ↪IMG_WIDTH]), label)
        ).map( # random horizontal flip
            lambda image, label: (tf.image.random_flip_left_right(image, seed=rnd.
        ↪seed(1234)), label), num_parallel_calls=num_threads
        ).map( # random translation
            lambda image, label: (tfa.image.translate(image, tf.random.
        ↪normal(shape=[2], seed=rnd.seed(1234))), label),
        ↪num_parallel_calls=num_threads
        ).repeat().batch(128).take(1).cache( # cache images
        ).prefetch(1)

test = test_ds.map(
    lambda image, label: (tf.image.convert_image_dtype(image, tf.float64),
    ↪label)
).map( # resize for input
    lambda image, label: (tf.image.resize(image, [IMG_HEIGHT, IMG_WIDTH]),
    ↪label)
).cache( # cache images
).map( # random horizontal flip
    lambda image, label: (tfa.image.translate(image, tf.random.
    ↪normal(shape=[2], seed=rnd.seed(1234))), label)
).map( # random translation
    lambda image, label: (tf.image.rgb_to_grayscale(image), label)
).shuffle(100).prefetch(1)

```

```

[76]: # Using pre trained MobileNetV2 model
# this model is trained on 224x224 images
mobnet = tf.keras.applications.MobileNetV2(include_top=False, weights=None)

```

```

# trainable parameters in model
nvar = len(mobnet.trainable_variables)
print('Number of variables: ', nvar)
for i in np.arange(nvar):
    print(mobnet.trainable_variables[i].name) # variable name
    print(mobnet.trainable_variables[i].shape) # kernel shape

```

```

Number of variables: 156
Conv1/kernel:0
(3, 3, 3, 32)
bn_Conv1/gamma:0
(32,)
bn_Conv1/beta:0
(32,)
expanded_conv_depthwise/depthwise_kernel:0
(3, 3, 32, 1)
expanded_conv_depthwise_BN/gamma:0
(32,)
expanded_conv_depthwise_BN/beta:0
(32,)
expanded_conv_project/kernel:0
(1, 1, 32, 16)
expanded_conv_project_BN/gamma:0
(16,)
expanded_conv_project_BN/beta:0
(16,)
block_1_expand/kernel:0
(1, 1, 16, 96)
block_1_expand_BN/gamma:0
(96,)
block_1_expand_BN/beta:0
(96,)
block_1_depthwise/depthwise_kernel:0
(3, 3, 96, 1)
block_1_depthwise_BN/gamma:0
(96,)
block_1_depthwise_BN/beta:0
(96,)
block_1_project/kernel:0
(1, 1, 96, 24)
block_1_project_BN/gamma:0
(24,)
block_1_project_BN/beta:0
(24,)
block_2_expand/kernel:0
(1, 1, 24, 144)

```

block_2_expand_BN/gamma:0
(144,)
block_2_expand_BN/beta:0
(144,)
block_2_depthwise/depthwise_kernel:0
(3, 3, 144, 1)
block_2_depthwise_BN/gamma:0
(144,)
block_2_depthwise_BN/beta:0
(144,)
block_2_project/kernel:0
(1, 1, 144, 24)
block_2_project_BN/gamma:0
(24,)
block_2_project_BN/beta:0
(24,)
block_3_expand/kernel:0
(1, 1, 24, 144)
block_3_expand_BN/gamma:0
(144,)
block_3_expand_BN/beta:0
(144,)
block_3_depthwise/depthwise_kernel:0
(3, 3, 144, 1)
block_3_depthwise_BN/gamma:0
(144,)
block_3_depthwise_BN/beta:0
(144,)
block_3_project/kernel:0
(1, 1, 144, 32)
block_3_project_BN/gamma:0
(32,)
block_3_project_BN/beta:0
(32,)
block_4_expand/kernel:0
(1, 1, 32, 192)
block_4_expand_BN/gamma:0
(192,)
block_4_expand_BN/beta:0
(192,)
block_4_depthwise/depthwise_kernel:0
(3, 3, 192, 1)
block_4_depthwise_BN/gamma:0
(192,)
block_4_depthwise_BN/beta:0
(192,)
block_4_project/kernel:0
(1, 1, 192, 32)

block_4_project_BN/gamma:0
(32,)
block_4_project_BN/beta:0
(32,)
block_5_expand/kernel:0
(1, 1, 32, 192)
block_5_expand_BN/gamma:0
(192,)
block_5_expand_BN/beta:0
(192,)
block_5_depthwise/depthwise_kernel:0
(3, 3, 192, 1)
block_5_depthwise_BN/gamma:0
(192,)
block_5_depthwise_BN/beta:0
(192,)
block_5_project/kernel:0
(1, 1, 192, 32)
block_5_project_BN/gamma:0
(32,)
block_5_project_BN/beta:0
(32,)
block_6_expand/kernel:0
(1, 1, 32, 192)
block_6_expand_BN/gamma:0
(192,)
block_6_expand_BN/beta:0
(192,)
block_6_depthwise/depthwise_kernel:0
(3, 3, 192, 1)
block_6_depthwise_BN/gamma:0
(192,)
block_6_depthwise_BN/beta:0
(192,)
block_6_project/kernel:0
(1, 1, 192, 64)
block_6_project_BN/gamma:0
(64,)
block_6_project_BN/beta:0
(64,)
block_7_expand/kernel:0
(1, 1, 64, 384)
block_7_expand_BN/gamma:0
(384,)
block_7_expand_BN/beta:0
(384,)
block_7_depthwise/depthwise_kernel:0
(3, 3, 384, 1)

block_7_depthwise_BN/gamma:0
(384,)
block_7_depthwise_BN/beta:0
(384,)
block_7_project/kernel:0
(1, 1, 384, 64)
block_7_project_BN/gamma:0
(64,)
block_7_project_BN/beta:0
(64,)
block_8_expand/kernel:0
(1, 1, 64, 384)
block_8_expand_BN/gamma:0
(384,)
block_8_expand_BN/beta:0
(384,)
block_8_depthwise/depthwise_kernel:0
(3, 3, 384, 1)
block_8_depthwise_BN/gamma:0
(384,)
block_8_depthwise_BN/beta:0
(384,)
block_8_project/kernel:0
(1, 1, 384, 64)
block_8_project_BN/gamma:0
(64,)
block_8_project_BN/beta:0
(64,)
block_9_expand/kernel:0
(1, 1, 64, 384)
block_9_expand_BN/gamma:0
(384,)
block_9_expand_BN/beta:0
(384,)
block_9_depthwise/depthwise_kernel:0
(3, 3, 384, 1)
block_9_depthwise_BN/gamma:0
(384,)
block_9_depthwise_BN/beta:0
(384,)
block_9_project/kernel:0
(1, 1, 384, 64)
block_9_project_BN/gamma:0
(64,)
block_9_project_BN/beta:0
(64,)
block_10_expand/kernel:0
(1, 1, 64, 384)

block_10_expand_BN/gamma:0
(384,)
block_10_expand_BN/beta:0
(384,)
block_10_depthwise/depthwise_kernel:0
(3, 3, 384, 1)
block_10_depthwise_BN/gamma:0
(384,)
block_10_depthwise_BN/beta:0
(384,)
block_10_project/kernel:0
(1, 1, 384, 96)
block_10_project_BN/gamma:0
(96,)
block_10_project_BN/beta:0
(96,)
block_11_expand/kernel:0
(1, 1, 96, 576)
block_11_expand_BN/gamma:0
(576,)
block_11_expand_BN/beta:0
(576,)
block_11_depthwise/depthwise_kernel:0
(3, 3, 576, 1)
block_11_depthwise_BN/gamma:0
(576,)
block_11_depthwise_BN/beta:0
(576,)
block_11_project/kernel:0
(1, 1, 576, 96)
block_11_project_BN/gamma:0
(96,)
block_11_project_BN/beta:0
(96,)
block_12_expand/kernel:0
(1, 1, 96, 576)
block_12_expand_BN/gamma:0
(576,)
block_12_expand_BN/beta:0
(576,)
block_12_depthwise/depthwise_kernel:0
(3, 3, 576, 1)
block_12_depthwise_BN/gamma:0
(576,)
block_12_depthwise_BN/beta:0
(576,)
block_12_project/kernel:0
(1, 1, 576, 96)

block_12_project_BN/gamma:0
(96,)
block_12_project_BN/beta:0
(96,)
block_13_expand/kernel:0
(1, 1, 96, 576)
block_13_expand_BN/gamma:0
(576,)
block_13_expand_BN/beta:0
(576,)
block_13_depthwise/depthwise_kernel:0
(3, 3, 576, 1)
block_13_depthwise_BN/gamma:0
(576,)
block_13_depthwise_BN/beta:0
(576,)
block_13_project/kernel:0
(1, 1, 576, 160)
block_13_project_BN/gamma:0
(160,)
block_13_project_BN/beta:0
(160,)
block_14_expand/kernel:0
(1, 1, 160, 960)
block_14_expand_BN/gamma:0
(960,)
block_14_expand_BN/beta:0
(960,)
block_14_depthwise/depthwise_kernel:0
(3, 3, 960, 1)
block_14_depthwise_BN/gamma:0
(960,)
block_14_depthwise_BN/beta:0
(960,)
block_14_project/kernel:0
(1, 1, 960, 160)
block_14_project_BN/gamma:0
(160,)
block_14_project_BN/beta:0
(160,)
block_15_expand/kernel:0
(1, 1, 160, 960)
block_15_expand_BN/gamma:0
(960,)
block_15_expand_BN/beta:0
(960,)
block_15_depthwise/depthwise_kernel:0
(3, 3, 960, 1)

```

block_15_depthwise_BN/gamma:0
(960,)
block_15_depthwise_BN/beta:0
(960,)
block_15_project/kernel:0
(1, 1, 960, 160)
block_15_project_BN/gamma:0
(160,)
block_15_project_BN/beta:0
(160,)
block_16_expand/kernel:0
(1, 1, 160, 960)
block_16_expand_BN/gamma:0
(960,)
block_16_expand_BN/beta:0
(960,)
block_16_depthwise/depthwise_kernel:0
(3, 3, 960, 1)
block_16_depthwise_BN/gamma:0
(960,)
block_16_depthwise_BN/beta:0
(960,)
block_16_project/kernel:0
(1, 1, 960, 320)
block_16_project_BN/gamma:0
(320,)
block_16_project_BN/beta:0
(320,)
Conv_1/kernel:0
(1, 1, 320, 1280)
Conv_1_bn/gamma:0
(1280,)
Conv_1_bn/beta:0
(1280,)

```

```

[77]: # clear keras session
      tf.keras.backend.clear_session()

```

0.1 Training from scratch

We are going to add a dropout layer and a new dense layer for 2 classes (i.e., cats and dogs)

```

[78]: dropout_rate = 0.2
      num_classes = 2
      CHANS = 3 # number of channels in input images
      input_shape = [IMG_HEIGHT, IMG_WIDTH, CHANS]

```

```

mbnet = tf.keras.applications.MobileNetV2(input_shape=input_shape,
    ↪include_top=False, weights=None)

model = tf.keras.Sequential([
    mbnet,
    tf.keras.layers.GlobalAveragePooling2D(),
    tf.keras.layers.Dropout(dropout_rate),
    tf.keras.layers.Dense(num_classes, activation='softmax')
], name='pre-trained_mobilenetV2')
model.build()

model.summary()

```

Model: "pre-trained_mobilenetV2"

| Layer (type) | Output Shape | Param # |
|---|--------------------|---------|
| mobilenetv2_1.00_224 (Functional) | (None, 7, 7, 1280) | 2257984 |
| global_average_pooling2d (GlobalAveragePooling2D) | (None, 1280) | 0 |
| dropout (Dropout) | (None, 1280) | 0 |
| dense (Dense) | (None, 2) | 2562 |

Total params: 2,260,546
 Trainable params: 2,226,434
 Non-trainable params: 34,112

```

[79]: with tf.GradientTape() as tape:
      tape.reset()

```

```

[81]: LR = 0.001 # learning rate
      optimizer = tf.optimizers.Adam(learning_rate=LR) # adam optimizer

      def train_step(model, X, Y):
          with tf.GradientTape() as tape:
              pred = model(X)
              current_loss = tf.reduce_mean(tf.losses.categorical_crossentropy(Y, pred))
              grads = tape.gradient(current_loss, model.trainable_variables)
              optimizer.apply_gradients(zip( grads , model.trainable_variables)) # update
          ↪gradients of all layers
              current_accuracy = tf.reduce_mean(tf.metrics.categorical_accuracy(Y, pred))

```

```
return(current_loss, current_accuracy)
```

```
[82]: import warnings
warnings.filterwarnings('ignore')
```

```
[83]: niter = 1000

tloss = []
tacc = []
vloss = []
vacc = []

for it in range(niter):
    for image, label in train:
        loss, acc = train_step( model , image , tf.one_hot(label, depth=2) )
    ↪#run training

    if it % 10 == 0: #log training metrics
        tf.print('iter: ',it, ', loss: {:.3f}, acc: {:.3f}'.format(loss, acc))
        tloss.append(loss)
        tacc.append(acc)

    if it % 50 == 0: #log validation metrics
        for val_image, val_label in valid:
            val_pred = model(val_image)
            val_loss = tf.reduce_mean(tf.losses.categorical_crossentropy(tf.
    ↪one_hot(val_label, depth=2), val_pred))
            val_acc = tf.reduce_mean(tf.metrics.categorical_accuracy(tf.
    ↪one_hot(val_label, depth=2), val_pred))
            tf.print('iter: ',it, ', validation loss: {:.3f}, validation acc: {:.
    ↪3f}'.format(val_loss, val_acc))
            vloss.append(val_loss)
            vacc.append(val_acc)
```

```
iter: 0 , loss: 0.693, acc: 0.438
iter: 0 , validation loss: 0.691, validation acc: 0.555
iter: 10 , loss: 0.693, acc: 0.531
iter: 20 , loss: 0.693, acc: 0.531
iter: 30 , loss: 0.693, acc: 0.500
iter: 40 , loss: 0.694, acc: 0.469
iter: 50 , loss: 0.691, acc: 0.594
iter: 50 , validation loss: 0.695, validation acc: 0.445
iter: 60 , loss: 0.693, acc: 0.500
iter: 70 , loss: 0.686, acc: 0.594
iter: 80 , loss: 0.689, acc: 0.562
iter: 90 , loss: 0.690, acc: 0.562
iter: 100 , loss: 0.690, acc: 0.562
```

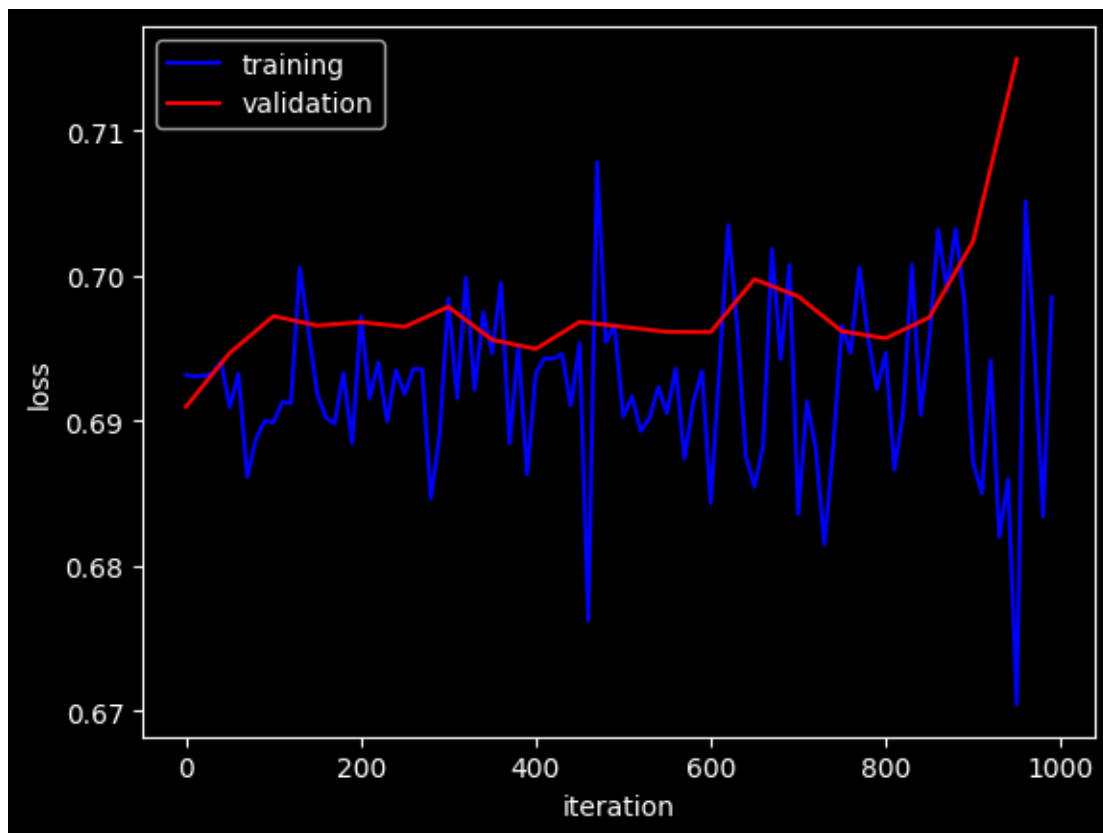
iter: 100 , validation loss: 0.697, validation acc: 0.445
iter: 110 , loss: 0.691, acc: 0.531
iter: 120 , loss: 0.691, acc: 0.531
iter: 130 , loss: 0.701, acc: 0.406
iter: 140 , loss: 0.696, acc: 0.438
iter: 150 , loss: 0.692, acc: 0.531
iter: 150 , validation loss: 0.697, validation acc: 0.445
iter: 160 , loss: 0.690, acc: 0.562
iter: 170 , loss: 0.690, acc: 0.594
iter: 180 , loss: 0.693, acc: 0.500
iter: 190 , loss: 0.688, acc: 0.625
iter: 200 , loss: 0.697, acc: 0.438
iter: 200 , validation loss: 0.697, validation acc: 0.445
iter: 210 , loss: 0.692, acc: 0.531
iter: 220 , loss: 0.694, acc: 0.500
iter: 230 , loss: 0.690, acc: 0.562
iter: 240 , loss: 0.693, acc: 0.500
iter: 250 , loss: 0.692, acc: 0.531
iter: 250 , validation loss: 0.696, validation acc: 0.445
iter: 260 , loss: 0.694, acc: 0.500
iter: 270 , loss: 0.694, acc: 0.500
iter: 280 , loss: 0.685, acc: 0.656
iter: 290 , loss: 0.689, acc: 0.562
iter: 300 , loss: 0.698, acc: 0.438
iter: 300 , validation loss: 0.698, validation acc: 0.445
iter: 310 , loss: 0.692, acc: 0.531
iter: 320 , loss: 0.700, acc: 0.375
iter: 330 , loss: 0.692, acc: 0.531
iter: 340 , loss: 0.697, acc: 0.406
iter: 350 , loss: 0.695, acc: 0.469
iter: 350 , validation loss: 0.696, validation acc: 0.445
iter: 360 , loss: 0.700, acc: 0.375
iter: 370 , loss: 0.688, acc: 0.594
iter: 380 , loss: 0.695, acc: 0.469
iter: 390 , loss: 0.686, acc: 0.656
iter: 400 , loss: 0.693, acc: 0.500
iter: 400 , validation loss: 0.695, validation acc: 0.445
iter: 410 , loss: 0.694, acc: 0.469
iter: 420 , loss: 0.694, acc: 0.469
iter: 430 , loss: 0.695, acc: 0.438
iter: 440 , loss: 0.691, acc: 0.562
iter: 450 , loss: 0.695, acc: 0.469
iter: 450 , validation loss: 0.697, validation acc: 0.445
iter: 460 , loss: 0.676, acc: 0.688
iter: 470 , loss: 0.708, acc: 0.344
iter: 480 , loss: 0.695, acc: 0.469
iter: 490 , loss: 0.697, acc: 0.438
iter: 500 , loss: 0.690, acc: 0.562

iter: 500 , validation loss: 0.696, validation acc: 0.445
iter: 510 , loss: 0.692, acc: 0.531
iter: 520 , loss: 0.689, acc: 0.562
iter: 530 , loss: 0.690, acc: 0.625
iter: 540 , loss: 0.692, acc: 0.562
iter: 550 , loss: 0.691, acc: 0.562
iter: 550 , validation loss: 0.696, validation acc: 0.445
iter: 560 , loss: 0.694, acc: 0.500
iter: 570 , loss: 0.687, acc: 0.625
iter: 580 , loss: 0.691, acc: 0.562
iter: 590 , loss: 0.693, acc: 0.500
iter: 600 , loss: 0.684, acc: 0.688
iter: 600 , validation loss: 0.696, validation acc: 0.445
iter: 610 , loss: 0.694, acc: 0.500
iter: 620 , loss: 0.703, acc: 0.344
iter: 630 , loss: 0.697, acc: 0.406
iter: 640 , loss: 0.688, acc: 0.594
iter: 650 , loss: 0.685, acc: 0.594
iter: 650 , validation loss: 0.700, validation acc: 0.445
iter: 660 , loss: 0.688, acc: 0.562
iter: 670 , loss: 0.702, acc: 0.438
iter: 680 , loss: 0.694, acc: 0.500
iter: 690 , loss: 0.701, acc: 0.438
iter: 700 , loss: 0.684, acc: 0.625
iter: 700 , validation loss: 0.699, validation acc: 0.445
iter: 710 , loss: 0.691, acc: 0.531
iter: 720 , loss: 0.688, acc: 0.562
iter: 730 , loss: 0.681, acc: 0.656
iter: 740 , loss: 0.688, acc: 0.594
iter: 750 , loss: 0.697, acc: 0.438
iter: 750 , validation loss: 0.696, validation acc: 0.445
iter: 760 , loss: 0.695, acc: 0.469
iter: 770 , loss: 0.701, acc: 0.312
iter: 780 , loss: 0.696, acc: 0.438
iter: 790 , loss: 0.692, acc: 0.531
iter: 800 , loss: 0.695, acc: 0.469
iter: 800 , validation loss: 0.696, validation acc: 0.445
iter: 810 , loss: 0.687, acc: 0.656
iter: 820 , loss: 0.691, acc: 0.562
iter: 830 , loss: 0.701, acc: 0.344
iter: 840 , loss: 0.690, acc: 0.562
iter: 850 , loss: 0.696, acc: 0.469
iter: 850 , validation loss: 0.697, validation acc: 0.445
iter: 860 , loss: 0.703, acc: 0.375
iter: 870 , loss: 0.699, acc: 0.438
iter: 880 , loss: 0.703, acc: 0.406
iter: 890 , loss: 0.698, acc: 0.469
iter: 900 , loss: 0.687, acc: 0.562


```
iter: 900 , validation loss: 0.702, validation acc: 0.445
iter: 910 , loss: 0.685, acc: 0.594
iter: 920 , loss: 0.694, acc: 0.500
iter: 930 , loss: 0.682, acc: 0.625
iter: 940 , loss: 0.686, acc: 0.562
iter: 950 , loss: 0.670, acc: 0.625
iter: 950 , validation loss: 0.715, validation acc: 0.445
iter: 960 , loss: 0.705, acc: 0.469
iter: 970 , loss: 0.695, acc: 0.406
iter: 980 , loss: 0.683, acc: 0.625
iter: 990 , loss: 0.698, acc: 0.438
```

```
[84]: titers = 10*np.arange(np.shape(tloss)[0])
      vitters = 50*np.arange(np.shape(vloss)[0])

      plt.plot(titers, tloss, c='b', label='training');
      plt.plot(vitters, vloss, c='r', label='validation');
      plt.legend();
      plt.xlabel('iteration');
      plt.ylabel('loss');
```




```
[85]: plt.plot(titers, tacc, c='b', label='training');
plt.plot(viters, vacc, c='r', label='validation');
plt.legend();
plt.xlabel('iteration');
plt.ylabel('accuracy');
```



0.2 Transfer learning

Now we will implement transfer learning from the imagenet weights

```
[86]: tf.keras.backend.clear_session() # clear keras session
```

```
[87]: dropout_rate = 0.2
num_classes = 2
IMG_HEIGHT = 224
IMG_WIDTH = 224
CHANS = 3
input_shape = [IMG_HEIGHT, IMG_WIDTH, CHANS]

mobnet = tf.keras.applications.MobileNetV2(input_shape=input_shape,
↳ include_top=False, weights='imagenet')
```

```

model = tf.keras.Sequential([
    mobnet,
    tf.keras.layers.GlobalAveragePooling2D(),
    tf.keras.layers.Dropout(dropout_rate),
    tf.keras.layers.Dense(num_classes, activation='softmax')
], name='MobileNetV2_weights_model')

model.build()

mobnet.trainable = False # freezes the first layers to the imagenet weights

model.summary()

```

Model: "MobileNetV2_weights_model"

| Layer (type) | Output Shape | Param # |
|---|--------------------|---------|
| mobilenetv2_1.00_224 (Functional) | (None, 7, 7, 1280) | 2257984 |
| global_average_pooling2d (GlobalAveragePooling2D) | (None, 1280) | 0 |
| dropout (Dropout) | (None, 1280) | 0 |
| dense (Dense) | (None, 2) | 2562 |

Total params: 2,260,546
 Trainable params: 2,562
 Non-trainable params: 2,257,984

```

[88]: with tf.GradientTape() as tape:
      tape.reset() # This resets gradient tape

```

```

[89]: LR = 1e-2 # learning rate
      optimizer = tf.keras.optimizers.Adam(LR) # initialize adam optimizer

      def train_step(model, X, Y):
          with tf.GradientTape() as tape:
              pred = model(X)
              current_loss = tf.reduce_mean(tf.losses.categorical_crossentropy(Y,
              ↪pred))
              grads = tape.gradient(current_loss, model.trainable_variables)
              optimizer.apply_gradients(zip(grads, model.trainable_variables))

```

```

current_accuracy = tf.reduce_mean(tf.metrics.categorical_accuracy(Y, pred))
return(current_loss, current_accuracy)

```

```

[90]: niter = 1000

tloss = []
tacc = []
vloss = []
vacc = []

for it in range(niter):
    for image, label in train:
        loss, acc = train_step(model , image , tf.one_hot(label,depth=2) ) #run_
        ↪training

    if it % 10 == 0: #log training metrics
        tf.print('iter: ',it, ', loss: {:.3f}, acc: {:.3f}'.format(loss, acc))
        tloss.append(loss)
        tacc.append(acc)

    if it % 50 == 0: #log validation metrics
        for val_image, val_label in valid:
            val_pred = model(val_image)
            val_loss = tf.reduce_mean(tf.losses.categorical_crossentropy(tf.
            ↪one_hot(val_label,depth=2) , val_pred))
            val_acc = tf.reduce_mean(tf.metrics.categorical_accuracy(tf.
            ↪one_hot(val_label,depth=2) , val_pred))
            tf.print('iter: ',it, ', validation loss: {:.3f}, validation acc: {:.
            ↪3f}'.format(val_loss, val_acc))
            vloss.append(val_loss)
            vacc.append(val_acc)

```

```

iter: 0 , loss: 0.839, acc: 0.438
iter: 0 , validation loss: 0.703, validation acc: 0.695
iter: 10 , loss: 0.054, acc: 0.969
iter: 20 , loss: 0.063, acc: 0.969
iter: 30 , loss: 0.086, acc: 0.938
iter: 40 , loss: 0.023, acc: 1.000
iter: 50 , loss: 0.000, acc: 1.000
iter: 50 , validation loss: 0.003, validation acc: 1.000
iter: 60 , loss: 0.002, acc: 1.000
iter: 70 , loss: 0.001, acc: 1.000
iter: 80 , loss: 0.029, acc: 1.000
iter: 90 , loss: 0.088, acc: 0.969
iter: 100 , loss: 0.006, acc: 1.000
iter: 100 , validation loss: 0.047, validation acc: 0.984
iter: 110 , loss: 0.022, acc: 1.000

```

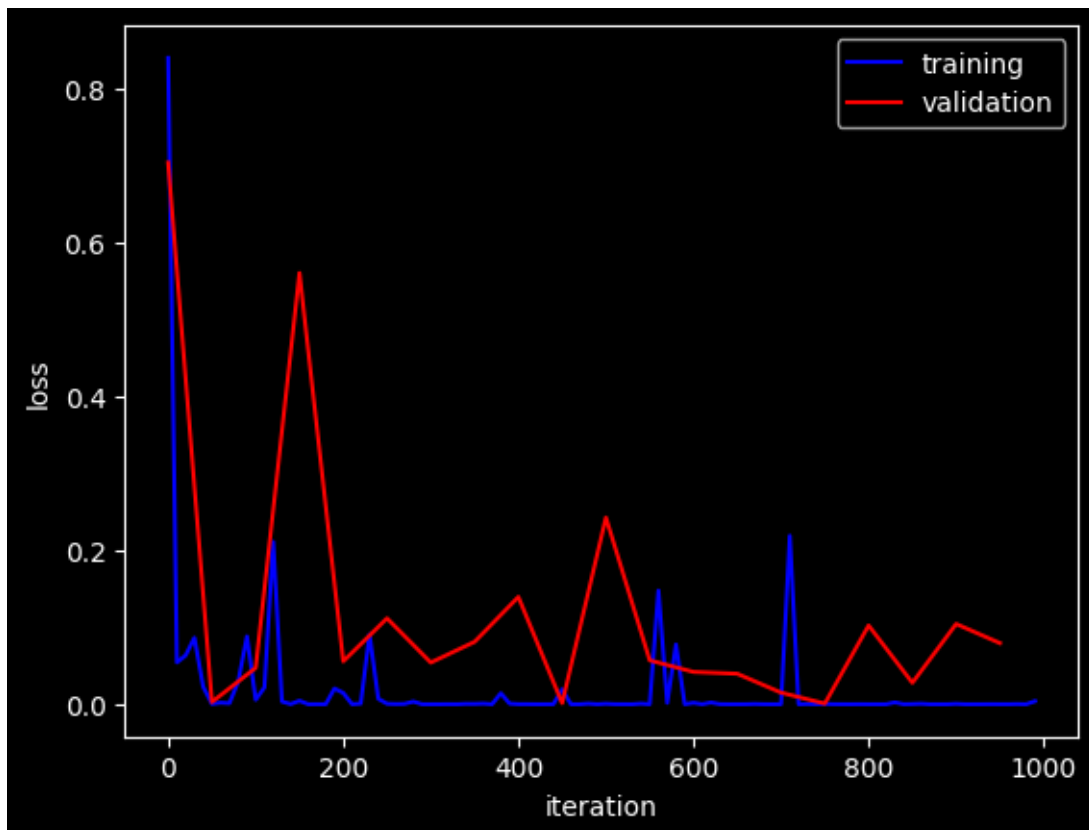
iter: 120 , loss: 0.211, acc: 0.938
iter: 130 , loss: 0.003, acc: 1.000
iter: 140 , loss: 0.000, acc: 1.000
iter: 150 , loss: 0.004, acc: 1.000
iter: 150 , validation loss: 0.560, validation acc: 0.914
iter: 160 , loss: 0.000, acc: 1.000
iter: 170 , loss: 0.000, acc: 1.000
iter: 180 , loss: 0.000, acc: 1.000
iter: 190 , loss: 0.020, acc: 1.000
iter: 200 , loss: 0.015, acc: 1.000
iter: 200 , validation loss: 0.056, validation acc: 0.984
iter: 210 , loss: 0.000, acc: 1.000
iter: 220 , loss: 0.001, acc: 1.000
iter: 230 , loss: 0.091, acc: 0.969
iter: 240 , loss: 0.007, acc: 1.000
iter: 250 , loss: 0.001, acc: 1.000
iter: 250 , validation loss: 0.111, validation acc: 0.984
iter: 260 , loss: 0.000, acc: 1.000
iter: 270 , loss: 0.000, acc: 1.000
iter: 280 , loss: 0.003, acc: 1.000
iter: 290 , loss: 0.000, acc: 1.000
iter: 300 , loss: 0.000, acc: 1.000
iter: 300 , validation loss: 0.054, validation acc: 0.984
iter: 310 , loss: 0.000, acc: 1.000
iter: 320 , loss: 0.000, acc: 1.000
iter: 330 , loss: 0.000, acc: 1.000
iter: 340 , loss: 0.001, acc: 1.000
iter: 350 , loss: 0.000, acc: 1.000
iter: 350 , validation loss: 0.081, validation acc: 0.984
iter: 360 , loss: 0.001, acc: 1.000
iter: 370 , loss: 0.000, acc: 1.000
iter: 380 , loss: 0.014, acc: 1.000
iter: 390 , loss: 0.001, acc: 1.000
iter: 400 , loss: 0.000, acc: 1.000
iter: 400 , validation loss: 0.139, validation acc: 0.969
iter: 410 , loss: 0.000, acc: 1.000
iter: 420 , loss: 0.000, acc: 1.000
iter: 430 , loss: 0.000, acc: 1.000
iter: 440 , loss: 0.000, acc: 1.000
iter: 450 , loss: 0.019, acc: 1.000
iter: 450 , validation loss: 0.001, validation acc: 1.000
iter: 460 , loss: 0.000, acc: 1.000
iter: 470 , loss: 0.000, acc: 1.000
iter: 480 , loss: 0.001, acc: 1.000
iter: 490 , loss: 0.000, acc: 1.000
iter: 500 , loss: 0.001, acc: 1.000
iter: 500 , validation loss: 0.242, validation acc: 0.961
iter: 510 , loss: 0.000, acc: 1.000

iter: 520 , loss: 0.000, acc: 1.000
iter: 530 , loss: 0.000, acc: 1.000
iter: 540 , loss: 0.001, acc: 1.000
iter: 550 , loss: 0.000, acc: 1.000
iter: 550 , validation loss: 0.057, validation acc: 0.992
iter: 560 , loss: 0.148, acc: 0.969
iter: 570 , loss: 0.002, acc: 1.000
iter: 580 , loss: 0.078, acc: 0.969
iter: 590 , loss: 0.000, acc: 1.000
iter: 600 , loss: 0.002, acc: 1.000
iter: 600 , validation loss: 0.042, validation acc: 0.969
iter: 610 , loss: 0.000, acc: 1.000
iter: 620 , loss: 0.002, acc: 1.000
iter: 630 , loss: 0.000, acc: 1.000
iter: 640 , loss: 0.000, acc: 1.000
iter: 650 , loss: 0.000, acc: 1.000
iter: 650 , validation loss: 0.040, validation acc: 0.984
iter: 660 , loss: 0.000, acc: 1.000
iter: 670 , loss: 0.000, acc: 1.000
iter: 680 , loss: 0.000, acc: 1.000
iter: 690 , loss: 0.000, acc: 1.000
iter: 700 , loss: 0.000, acc: 1.000
iter: 700 , validation loss: 0.015, validation acc: 0.992
iter: 710 , loss: 0.219, acc: 0.969
iter: 720 , loss: 0.000, acc: 1.000
iter: 730 , loss: 0.000, acc: 1.000
iter: 740 , loss: 0.000, acc: 1.000
iter: 750 , loss: 0.000, acc: 1.000
iter: 750 , validation loss: 0.001, validation acc: 1.000
iter: 760 , loss: 0.000, acc: 1.000
iter: 770 , loss: 0.000, acc: 1.000
iter: 780 , loss: 0.000, acc: 1.000
iter: 790 , loss: 0.000, acc: 1.000
iter: 800 , loss: 0.000, acc: 1.000
iter: 800 , validation loss: 0.102, validation acc: 0.977
iter: 810 , loss: 0.000, acc: 1.000
iter: 820 , loss: 0.000, acc: 1.000
iter: 830 , loss: 0.002, acc: 1.000
iter: 840 , loss: 0.000, acc: 1.000
iter: 850 , loss: 0.000, acc: 1.000
iter: 850 , validation loss: 0.028, validation acc: 0.992
iter: 860 , loss: 0.001, acc: 1.000
iter: 870 , loss: 0.000, acc: 1.000
iter: 880 , loss: 0.000, acc: 1.000
iter: 890 , loss: 0.000, acc: 1.000
iter: 900 , loss: 0.001, acc: 1.000
iter: 900 , validation loss: 0.104, validation acc: 0.992
iter: 910 , loss: 0.000, acc: 1.000

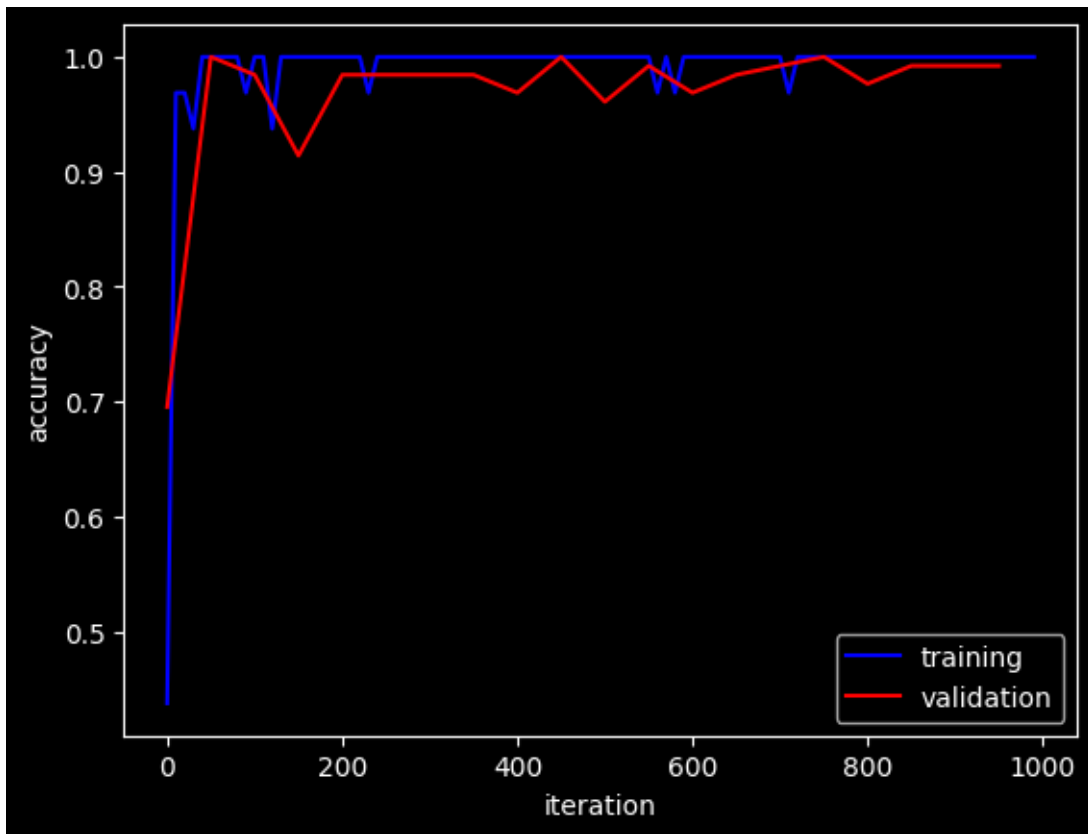
```
iter: 920 , loss: 0.000, acc: 1.000
iter: 930 , loss: 0.000, acc: 1.000
iter: 940 , loss: 0.000, acc: 1.000
iter: 950 , loss: 0.000, acc: 1.000
iter: 950 , validation loss: 0.079, validation acc: 0.992
iter: 960 , loss: 0.000, acc: 1.000
iter: 970 , loss: 0.000, acc: 1.000
iter: 980 , loss: 0.000, acc: 1.000
iter: 990 , loss: 0.004, acc: 1.000
```

```
[91]: # plot training and validation loss as a function of iteration
titors = 10*np.arange(np.shape(tloss)[0])
vitors = 50*np.arange(np.shape(vloss)[0])

plt.plot(titors, tloss, c='b', label='training');
plt.plot(vitors, vloss, c='r', label='validation');
plt.legend();
plt.xlabel('iteration');
plt.ylabel('loss');
```



```
[92]: # plot training and validation accuracy as a function of iteration
plt.plot(titers, tacc, c='b', label='training');
plt.plot(viters, vacc, c='r', label='validation');
plt.legend();
plt.xlabel('iteration');
plt.ylabel('accuracy');
```



0.3 Save weights

```
[93]: # Lets save the trained weights into hdf5 format
import os
CODE_PATH = '/Users/rahuln/vscode/'
MODEL_PATH = os.path.join(CODE_PATH, 'mlis_workshops', 'workshop4', 'my_weights.
↳hdf5')
model.save_weights(MODEL_PATH)
```

```
[94]: tf.keras.backend.clear_session()
```

0.4 Load weights

```
[95]: # Load previously saved weights onto new model

dropoutrate = 0.2
num_classes = 2
input_shape = [IMG_HEIGHT, IMG_WIDTH, CHANS]

mobnet = tf.keras.applications.MobileNetV2(input_shape=input_shape,
    ↪include_top=False, weights='imagenet') #We now dont want randomized weights
    ↪but to load weights from imagenet

model = tf.keras.Sequential([
    mobnet,
    tf.keras.layers.GlobalAveragePooling2D(),
    tf.keras.layers.Dropout(dropoutrate),
    tf.keras.layers.Dense(num_classes, activation='softmax')
], name='loaded_sequential_model')
model.build()

mobnet.trainable = False # freeze the first layers to the imagenet weights

model.summary() # print the model
```

Model: "loaded_sequential_model"

| Layer (type) | Output Shape | Param # |
|---|--------------------|---------|
| mobilenetv2_1.00_224 (Functional) | (None, 7, 7, 1280) | 2257984 |
| global_average_pooling2d (GlobalAveragePooling2D) | (None, 1280) | 0 |
| dropout (Dropout) | (None, 1280) | 0 |
| dense (Dense) | (None, 2) | 2562 |

=====
Total params: 2,260,546
Trainable params: 2,562
Non-trainable params: 2,257,984
=====

```
[96]: # Fine tuning the model by training with a smaller learning rate

LR = 1e-5 # learning rate
```



```

optimizer = tf.optimizers.Adam(LR) # adam optimizer

def train_step( model, X , Y):
    with tf.GradientTape() as tape:
        pred = model( X )
        current_loss = tf.reduce_mean(tf.losses.categorical_crossentropy( Y, 
↳pred))
        grads = tape.gradient(current_loss, model.trainable_variables)
        optimizer.apply_gradients( zip( grads , model.trainable_variables) )
        current_accuracy = tf.reduce_mean(tf.metrics.categorical_accuracy(Y, pred))
    return(current_loss, current_accuracy)

```

```

[97]: niter = 1000

tloss = []
tacc = []
vloss = []
vacc = []

for it in range(niter):
    for image, label in train:
        loss, acc = train_step( model , image , tf.one_hot(label,depth=2) )↳
↳#run training

    if it % 10 is 0: #log training metrics
        tf.print('iter: ',it, ', loss: {:.3f}, acc: {:.3f}'.format(loss, acc))
        tloss.append(loss)
        tacc.append(acc)

    if it % 50 is 0: #log validation metrics
        for val_image, val_label in valid:
            val_pred = model(val_image)
            val_loss = tf.reduce_mean(tf.losses.categorical_crossentropy(tf.
↳one_hot(val_label,depth=2) , val_pred))
            val_acc = tf.reduce_mean(tf.metrics.categorical_accuracy(tf.
↳one_hot(val_label,depth=2) , val_pred))
            tf.print('iter: ',it, ', validation loss: {:.3f}, validation acc: {:.
↳3f}'.format(val_loss, val_acc))
            vloss.append(val_loss)
            vacc.append(val_acc)

```

```

iter: 0 , loss: 0.889, acc: 0.438
iter: 0 , validation loss: 0.794, validation acc: 0.500
iter: 10 , loss: 0.721, acc: 0.562
iter: 20 , loss: 0.710, acc: 0.594
iter: 30 , loss: 0.650, acc: 0.625
iter: 40 , loss: 0.869, acc: 0.438

```

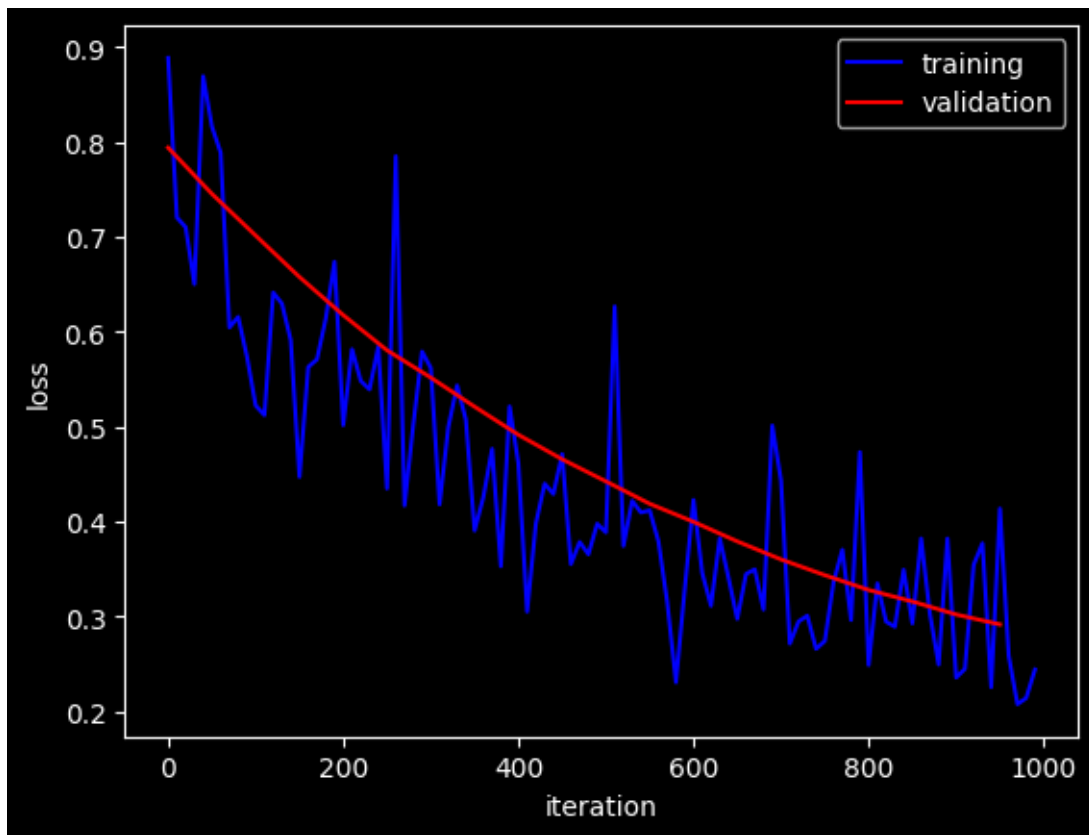
iter: 50 , loss: 0.816, acc: 0.562
iter: 50 , validation loss: 0.745, validation acc: 0.531
iter: 60 , loss: 0.789, acc: 0.500
iter: 70 , loss: 0.605, acc: 0.656
iter: 80 , loss: 0.616, acc: 0.719
iter: 90 , loss: 0.574, acc: 0.750
iter: 100 , loss: 0.522, acc: 0.750
iter: 100 , validation loss: 0.701, validation acc: 0.617
iter: 110 , loss: 0.512, acc: 0.719
iter: 120 , loss: 0.641, acc: 0.625
iter: 130 , loss: 0.630, acc: 0.531
iter: 140 , loss: 0.591, acc: 0.594
iter: 150 , loss: 0.447, acc: 0.812
iter: 150 , validation loss: 0.658, validation acc: 0.641
iter: 160 , loss: 0.563, acc: 0.688
iter: 170 , loss: 0.571, acc: 0.688
iter: 180 , loss: 0.614, acc: 0.719
iter: 190 , loss: 0.674, acc: 0.625
iter: 200 , loss: 0.502, acc: 0.781
iter: 200 , validation loss: 0.618, validation acc: 0.680
iter: 210 , loss: 0.582, acc: 0.688
iter: 220 , loss: 0.548, acc: 0.750
iter: 230 , loss: 0.539, acc: 0.625
iter: 240 , loss: 0.584, acc: 0.656
iter: 250 , loss: 0.435, acc: 0.812
iter: 250 , validation loss: 0.581, validation acc: 0.711
iter: 260 , loss: 0.785, acc: 0.625
iter: 270 , loss: 0.417, acc: 0.844
iter: 280 , loss: 0.502, acc: 0.812
iter: 290 , loss: 0.579, acc: 0.625
iter: 300 , loss: 0.562, acc: 0.750
iter: 300 , validation loss: 0.552, validation acc: 0.742
iter: 310 , loss: 0.418, acc: 0.906
iter: 320 , loss: 0.500, acc: 0.812
iter: 330 , loss: 0.544, acc: 0.719
iter: 340 , loss: 0.508, acc: 0.781
iter: 350 , loss: 0.391, acc: 0.906
iter: 350 , validation loss: 0.521, validation acc: 0.742
iter: 360 , loss: 0.426, acc: 0.844
iter: 370 , loss: 0.477, acc: 0.875
iter: 380 , loss: 0.353, acc: 0.906
iter: 390 , loss: 0.521, acc: 0.781
iter: 400 , loss: 0.461, acc: 0.844
iter: 400 , validation loss: 0.492, validation acc: 0.781
iter: 410 , loss: 0.305, acc: 0.938
iter: 420 , loss: 0.398, acc: 0.906
iter: 430 , loss: 0.440, acc: 0.875
iter: 440 , loss: 0.429, acc: 0.875

iter: 450 , loss: 0.471, acc: 0.781
iter: 450 , validation loss: 0.466, validation acc: 0.797
iter: 460 , loss: 0.355, acc: 0.969
iter: 470 , loss: 0.379, acc: 0.906
iter: 480 , loss: 0.366, acc: 0.938
iter: 490 , loss: 0.398, acc: 0.844
iter: 500 , loss: 0.389, acc: 0.906
iter: 500 , validation loss: 0.442, validation acc: 0.820
iter: 510 , loss: 0.627, acc: 0.750
iter: 520 , loss: 0.374, acc: 0.875
iter: 530 , loss: 0.422, acc: 0.812
iter: 540 , loss: 0.410, acc: 0.844
iter: 550 , loss: 0.412, acc: 0.906
iter: 550 , validation loss: 0.419, validation acc: 0.852
iter: 560 , loss: 0.379, acc: 0.875
iter: 570 , loss: 0.315, acc: 0.969
iter: 580 , loss: 0.231, acc: 0.969
iter: 590 , loss: 0.329, acc: 0.938
iter: 600 , loss: 0.423, acc: 0.781
iter: 600 , validation loss: 0.400, validation acc: 0.852
iter: 610 , loss: 0.346, acc: 0.906
iter: 620 , loss: 0.311, acc: 0.906
iter: 630 , loss: 0.383, acc: 0.875
iter: 640 , loss: 0.341, acc: 0.875
iter: 650 , loss: 0.298, acc: 0.906
iter: 650 , validation loss: 0.379, validation acc: 0.867
iter: 660 , loss: 0.345, acc: 0.844
iter: 670 , loss: 0.350, acc: 0.906
iter: 680 , loss: 0.307, acc: 0.938
iter: 690 , loss: 0.501, acc: 0.812
iter: 700 , loss: 0.443, acc: 0.875
iter: 700 , validation loss: 0.360, validation acc: 0.875
iter: 710 , loss: 0.271, acc: 0.938
iter: 720 , loss: 0.295, acc: 0.938
iter: 730 , loss: 0.301, acc: 0.938
iter: 740 , loss: 0.266, acc: 0.906
iter: 750 , loss: 0.274, acc: 0.938
iter: 750 , validation loss: 0.344, validation acc: 0.883
iter: 760 , loss: 0.336, acc: 0.938
iter: 770 , loss: 0.371, acc: 0.875
iter: 780 , loss: 0.296, acc: 0.938
iter: 790 , loss: 0.473, acc: 0.844
iter: 800 , loss: 0.249, acc: 0.906
iter: 800 , validation loss: 0.328, validation acc: 0.883
iter: 810 , loss: 0.335, acc: 0.875
iter: 820 , loss: 0.295, acc: 0.938
iter: 830 , loss: 0.289, acc: 0.938
iter: 840 , loss: 0.349, acc: 0.969

```
iter: 850 , loss: 0.293, acc: 0.938
iter: 850 , validation loss: 0.316, validation acc: 0.883
iter: 860 , loss: 0.382, acc: 0.906
iter: 870 , loss: 0.303, acc: 0.906
iter: 880 , loss: 0.249, acc: 0.938
iter: 890 , loss: 0.382, acc: 0.875
iter: 900 , loss: 0.236, acc: 0.969
iter: 900 , validation loss: 0.302, validation acc: 0.883
iter: 910 , loss: 0.245, acc: 0.938
iter: 920 , loss: 0.354, acc: 0.844
iter: 930 , loss: 0.377, acc: 0.906
iter: 940 , loss: 0.226, acc: 0.969
iter: 950 , loss: 0.414, acc: 0.781
iter: 950 , validation loss: 0.292, validation acc: 0.883
iter: 960 , loss: 0.259, acc: 0.938
iter: 970 , loss: 0.208, acc: 0.906
iter: 980 , loss: 0.214, acc: 0.969
iter: 990 , loss: 0.244, acc: 0.938
```

```
[98]: # plot training and validation loss as a function of iteration
titors = 10*np.arange(np.shape(tloss)[0])
vitors = 50*np.arange(np.shape(vloss)[0])

plt.plot(titors, tloss, c='b', label='training');
plt.plot(vitors, vloss, c='r', label='validation');
plt.legend();
plt.xlabel('iteration');
plt.ylabel('loss');
```



```
[99]: plt.plot(titers, tacc, c='b', label='training');  
plt.plot(viters, vacc, c='r', label='validation');  
plt.legend();  
plt.xlabel('iteration');  
plt.ylabel('accuracy');
```

